

PennyLane Search Exploratory Data Analysis

We will be exploring a dataset regarding the recent searches from PennyLane.ai/search. We will start by loading the data and getting a general idea of what kinds of data are including in the dataset.

```
In [1]: #import required libraries
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

#read in csv file of search data
df=pd.read_csv('/Users/lindsaybabcock/Documents/Work/Applications/Xanadu')
```

```
In [2]: #Change column to datetime types
df['created_at']=pd.to_datetime(df['created_at'],
                                format='%Y-%m-%d %H:%M:%S.%f')

#take a look at data types/data
df.info()
df.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 220524 entries, 0 to 220523
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     220524 non-null object
1   created_at                           220524 non-null datetime64[ns, UTC]
2   query_text                           19596 non-null  object
3   total_results                        220524 non-null int64
4   search_parameters                    220524 non-null object
5   query_content_types                  220456 non-null object
6   response_content_ids                 220524 non-null object
7   role                                 576 non-null    object
8   survey                              3095 non-null   object
9   personas                            3095 non-null   object
10  persona_other                        112 non-null    object
dtypes: datetime64[ns, UTC](1), int64(1), object(9)
memory usage: 18.5+ MB
```

```
Out[2]:
```

	id	created_at	query_text	total_results	search_parameters	query_con
0	04c0a197-0cbc-4bca-b5d4-c6134ba49546	2024-08-08 21:25:44.101779+00:00	mikhail	4	{"sort": null, "limit": 16, "filter": "((type ...	
1	0b0369f3-2bc9-4aa2-9f92-27e8954d8c41	2024-08-08 21:54:19.770839+00:00	NaN	37	{"sort": ["published_at:desc"], "limit": 16, "...	
2	0d72dae5-3aac-469f-910b-a7b98597ef7d	2024-08-08 21:26:07.410746+00:00	mikhail	5	{"sort": null, "limit": 16, "filter": "((type ...	
3	0e4fb2dd-6d5d-449b-b6a3-374d5430534a	2024-08-08 21:26:50.247836+00:00	izaac	82	{"sort": null, "limit": 16, "filter": "((type ...	
4	17d8319b-360a-4bf5-84a6-f2b1c753a692	2024-08-08 22:07:49.601295+00:00	NaN	5343	{"sort": null, "limit": 16, "filter": "((type ...	

Search Traffic

Now that we have an idea of the data we are exploring, we will begin by looking at trends in the search traffic over time.


```
In [3]: #Aggregate data to show searches per day
df['created_at'].dt.date.value_counts().plot(rot=25, title='PennyLane Q
plt.axhline(df['created_at'].dt.date.value_counts().mean(),c='r',ls='--
plt.ylabel('Number of Queries')
plt.xlabel('Date')
plt.legend(['query count', 'mean query count']);

print('Searches per day:')
print(round(df['created_at'].dt.date.value_counts().describe()))
```

Searches per day:

count	79.0
mean	2791.0
std	4186.0
min	58.0
25%	590.0
50%	755.0
75%	3546.0
max	18935.0

Name: created_at, dtype: float64



```
In [4]: #Show top queries during the time of the search traffic spike.
df_peak = df.loc[(df['created_at'] >= '2024-08-27') & (df['created_at']
print('Top searches during peak period:')
print(df_peak['query_text'].value_counts(dropna=False).rename_axis('Que

#Show top queries outside of the time of the search traffic spike.
df_nonpeak = df.loc[(df['created_at'] < '2024-08-27') | (df['created_at']
print('Top searches outside of peak period:')
print(df_nonpeak['query_text'].value_counts(dropna=False).rename_axis('
```

Top searches during peak period:

	Query	Count
0	NaN	174394
1	test	3010
2	vqe	71
3	qaoa	66
4	neural network	37

Top searches outside of peak period:

	Query	Count
0	NaN	26534
1	test	4622
2	quantum	192
3	qaoa	178
4	vqe	143

```
In [5]: #create subset dataframe excluding nulls in the query text
df_text = df.dropna(subset=['query_text'])

#Aggregate data to show searches per day of non-null queries
df_text['created_at'].dt.date.value_counts().plot(rot=25, title='Non-null queries',
plt.axhline(df_text['created_at'].dt.date.value_counts().mean(),c='r',ls='dashed')
plt.ylabel('Number of Queries')
plt.xlabel('Date')
plt.legend(['query count', 'mean query count']);


print('Non-null searches per day:')
print(round(df_text['created_at'].dt.date.value_counts().describe()))
print('Total non-null searches: ', df_text.shape[0])
```

Non-null searches per day:

```
count      79.0
mean       248.0
std         74.0
min         36.0
25%        196.0
50%        252.0
75%        293.0
max         482.0
```

Name: created_at, dtype: float64

Total non-null searches: 19596



There was a significant increase in search traffic between late August to mid September.

Looking more closely at what was being queried at this time, it appears searches containing text did not increase in the same way. I am unsure of what caused the spike in activity of null queries during this time, potentially site testing/development.

Removing the null queries brings the average queries per day down drastically from 2791 to 248. For the remainder of the analysis we will focus on queries containing text only.

```
In [7]: #import warnings
#warnings.filterwarnings('ignore')

df_text['day_of_week']=df_text['created_at'].dt.dayofweek
df_text['hour_of_day']=df_text['created_at'].dt.hour

search_counts = df_text.groupby(['day_of_week', 'hour_of_day']).size().reset_index()
search_counts.index = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']

plt.figure(figsize=(12, 8))
sns.heatmap(search_counts, cbar_kws={'label': 'Total queries per Time Unit'})
plt.title('Frequency of Queries by Day of Week and Hour of Day')
plt.xlabel('Hour of Day')
plt.ylabel('Day of Week')
plt.show()
```



I will be assuming that the 'created_at' times are recorded in the local time zone of the user. Here we see that the majority of queries are made during normal working hours - Monday to Friday, midday, suggesting that users of PennyLane are using it for work/school purposes.

Demographics

Below we will explore the distribution of demographics amongst users making queries on PennyLane.ai/search.

```
In [8]: #Aggregate data to find percentage of users that log-in and use pennylane
df_text['survey'].value_counts(dropna=False, normalize=True).mul(100).re
```


```
Out[8]: NaN          96.08%
{"uses_pennylane": true}    2.92%
{"uses_pennylane": false}   1.0%
Name: survey, dtype: object
```

96% of queries are done by users that are not signed in, thus we have little data about their specific roles. It may be useful to include features to encourage users to sign in/make accounts.

```
In [9]: #Show distribution of personas
#print(df_text['personas'].unique())
persona_types = ['RESEARCHER', 'DEVELOPER', 'STUDENT', 'HOBBYIST', 'EDUCATOR']
df_personas = pd.DataFrame(persona_types, columns=['type'])
df_personas['counts'] = [df_text['personas'].str.contains(x).sum() for x in persona_types]
df_personas.plot(kind='bar', x='type', legend=False)
plt.xlabel('Number of PennyLane Queries by Persona Type')
plt.ylabel('Number of Queries')

#Show non-categorized personas
print(df_text['persona_other'].value_counts().rename_axis('persona_other'))
```

	persona_other	Count
0	Community manager	28
1	Xanadu employee	5
2	Senior Product Manager	3



```

In [10]: #create subset dataframes for various personas
df_researcher = df[df.personas.str.contains('RESEARCHER',na=False)]
df_developer = df[df.personas.str.contains('DEVELOPER',na=False)]
df_student = df[df.personas.str.contains('STUDENT',na=False)]
df_hobbyist = df[df.personas.str.contains('HOBBYIST',na=False)]
df_educator = df[df.personas.str.contains('EDUCATOR',na=False)]
#df_other = df[df.personas.str.contains("OTHER",na=False)] #choosing no

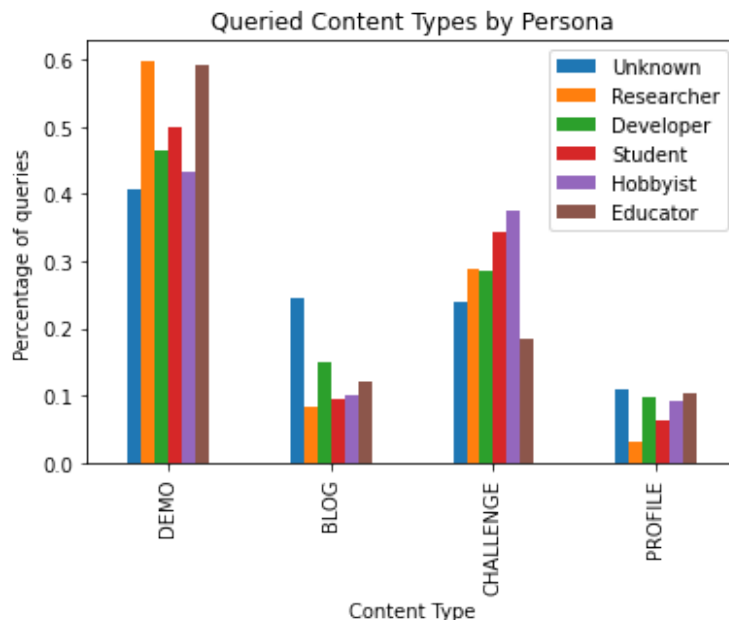
content_types = ['DEMO','BLOG','CHALLENGE','PROFILE']
personas_df = [df_text,df_researcher,df_developer,df_student,df_hobbyist]

#count occurrences of queries filtered for each content type for each persona
df_gen = pd.DataFrame(content_types,columns=['type'])
for dataframe in range(0,len(personas_df)):
    df_gen[dataframe] = [personas_df[dataframe]['query_content_types'].str.count(content_types)]
    df_gen[dataframe] = df_gen[dataframe]/df_gen[dataframe].sum()

#Show barplot of content types for each persona
df_gen.plot(kind='bar', x='type')
plt.legend(['Unknown','Researcher','Developer','Student','Hobbyist','Educator'])
plt.title('Queried Content Types by Persona')
plt.xlabel('Content Type')
plt.ylabel('Percentage of queries')

```

Out[10]: Text(0, 0.5, 'Percentage of queries')



The highest number of queries are made by Students, Researchers, and Developers. If we focus on supporting Researchers, as they are a high value persona, we can see that researchers are most often looking for demos, roughly 60% of the time, when searching PennyLane.ai. Focusing on providing high quality demo webpages would likely improve user experience for Researchers. On the other hand, researchers appear to be uninterested in profile or blog posts, searching for those types of content less than 10% of the time. For all groups, demos are most commonly searched for, therefore it may be beneficial to return them more often or closer to the top of the page for general queries without these filters on.

Top Queries

We can now explore what is being queried to understand the users top interests.

```
In [11]: #Count which text items were queried the most
print('General top queries:')
print(df_text['query_text'].value_counts().rename_axis('query_text').reset_index())

print('Top queries of researchers:')
print(df_researcher['query_text'].value_counts().rename_axis('query_text').reset_index())
```

General top queries:

	query_text	Count
0	test	7632
1	qaoa	244
2	quantum	228
3	vqe	214
4	qml	149
5	quantum neural network	113
6	variational	109
7	quantum machine learning	101
8	torch	84
9	hamiltonian	81
10	neural network	80
11	shor	78
12	gpu	69
13	classical shadow	68
14	kernel	68

Top queries of researchers:

	query_text	Count
0	simulated annealing	12
1	qaoa	11
2	hamiltonian	9
3	phase kickback	8
4	qmldata	8
5	optimization	7
6	quantum machine learning	7
7	gpu	7
8	photonic	6
9	encoding	6
10	kottmann	6
11	vqe	5
12	grover	5
13	certificate	5
14	image	5

Looking at the top searches (ignoring 'test'), we can see that most often users are searching for items such as 'Quantum approximate optimization algorithm (QAOA)', 'Quantum Machine Learning (QML)', and 'Hamiltonian'. Researchers are also often searching for those items, as well as 'Simulated Annealing', 'Phase Kick-back', and 'Optimization'.

Next we can order query results by the total results returned for the query in ascending order, and the number of times researchers have searched them, in descending order. These results give us an idea of what concepts users are interested in that are lacking resources on the

PennyLane.ai website.

```
In [12]: #Show lowest
df_totals=df_text.groupby(['query_text','total_results']).size().reset_
df_totals=df_totals.sort_values(['total_results', 'count'], ascending =
print('Low result queries:')
print(df_totals.loc[df_totals['count'] > 1].head(20)) #exclude entries w

print('Low result queries for researchers:')
df_totals_research=df_researcher.groupby(['query_text','total_results'])
df_totals_research=df_totals_research.sort_values(['total_results', 'co
print(df_totals_research.loc[df_totals_research['count'] > 1].head(15))
```

Low result queries:

	query_text	total_results	count
77	alan	0	19
5486	vqc	0	19
338	bb84	0	10
447	building molecular hamiltonians	0	9
897	cryptography	0	8
2881	mqt	0	8
4421	qubo	0	7
101	alpine	0	6
261	autoencoder	0	6
3568	qkd	0	6
600	cite	0	5
3726	qpc	0	5
4584	rl	0	5
5278	tsp	0	5
1414	federated	0	4
1656	gpu	0	4
1940	hhl	0	4
2917	nasa	0	4
3343	poissonboltzmann	0	4
3455	qaoa	0	4

Low result queries for researchers:

	query_text	total_results	count
182	womanium challenge	0	3
33	diamond	0	2
40	gpu	0	2
69	iqc	0	2
114	qaqo	0	2
127	quantum bit string comparator	0	2
82	matt silverman	1	2
105	photonic	1	2
178	vqd	1	2
41	gpu	2	2
149	simulated annealing	2	2
108	photonics	3	2
176	vqa	5	4
121	qnn	5	2
42	gpu	6	2

In general, there appear to be a fair number of searches for quantum cryptography related queries ('bb84','cryptography','qkd'), with zero returned results. Specifically for researchers there are more quantum hardware related queries such as 'photonics' or 'diamond'. It may be helpful for users if more resources on these topics were developed.

Further, certain results suggest user typos, such as 'qaqo', where the user is likely looking for 'qaoa'. Indicating potential room for improvement in the software for typo suggestion.

As with many datasets, further exploratory analysis could be done, including but not limited to: expanding the top queries analysis to include the users filters on content type (ie. see if diamond+BLOG is more often queried than diamond+DEMO), as well as explore sorting filters on queries, and understanding which result content id are shown the most.

In []: